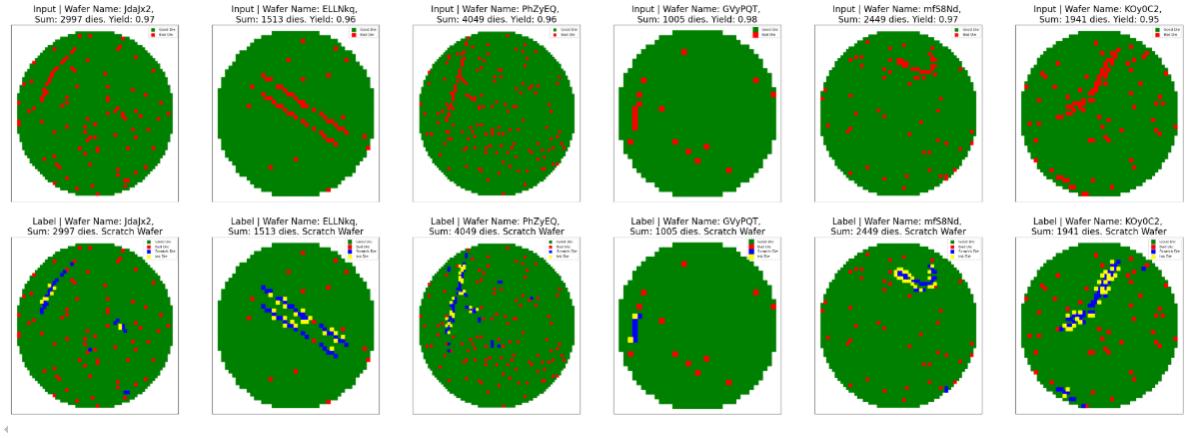


On the first page, I'll show the results, and then the process.

Results:

This is the result of my model on the test set:



As a business consideration, I added an output indicating whether the wafer is scratched and the size of the scratch:

	WaferName	HasScratch	ScratchPercent
0	ELLNkq	True	3.899537
1	GVyPQT	True	0.796020
2	JdaJx2	True	1.067734
3	KOyOC2	True	3.967027
4	PhZyEQ	True	1.951099
5	mfS8Nd	True	2.000817

The process I went through:

Found different wafer sizes. Checked if size affects defect rate. No strong link, but smallest wafers seem slightly more prone to defects.

Average defect rate by wafer size:

	WaferSize	DefectRate
0	1005.0	0.041189
1	1513.0	0.037805
2	1941.0	0.038410
3	2449.0	0.036768
4	2997.0	0.035885
5	3405.0	0.035086
6	4049.0	0.034905

Features it creates:

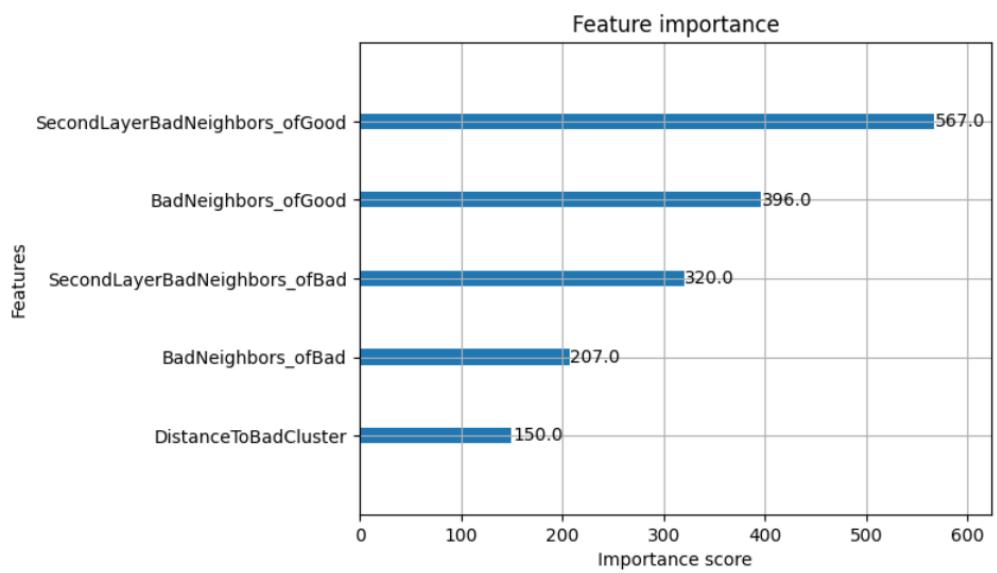
BadNeighbors_ofGood: Number of "bad" neighbors (`IsGoodDie == False`) within a distance of 1.5 around good dies.

BadNeighbors_ofBad: Same as above, but around bad dies.

SecondLayerBadNeighbors_ofGood: Number of bad dies in the second layer (neighbors of neighbors) around good dies.

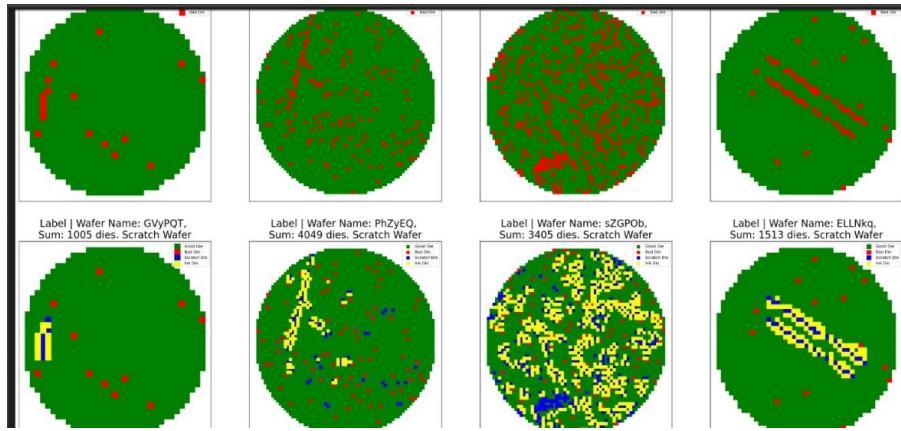
SecondLayerBadNeighbors_ofBad: Same as above, for bad dies.

DistanceToBadCluster: Distance from each die to the nearest bad die.



Model

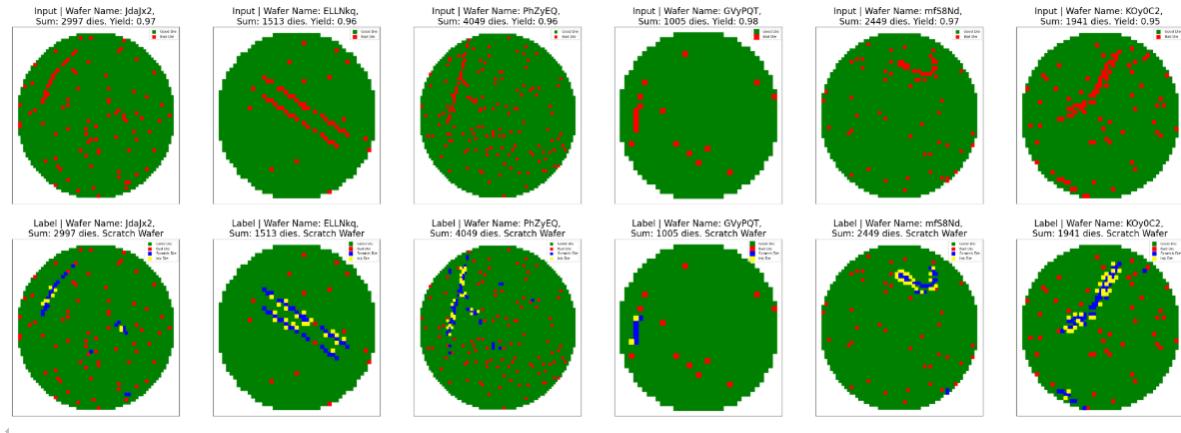
Tested XGBoost and Random Forest—XGBoost performed better.
Chose simple, efficient models over neural networks for faster, low-resource deployment, aligning with business constraints.



At first, the model tended to detect scratches as much wider than they actually were.

I preferred using XGBoost because it allows fine-tuning the scale_pos_weight parameter, which is especially helpful in this highly imbalanced dataset—where good dies significantly outnumber bad ones.

After some tuning, I achieved:



Despite a few misses, the model successfully detects scratches.
With **85% recall**, it's effective in identifying most defective dies:

Accuracy:	0.9903
Precision (True class):	0.4439
Recall (True class):	0.8500
F1-score (True class):	0.5832
ROC AUC:	0.9938
Classification Report:	
	precision recall f1-score support
False	1.00 0.99 1.00 361632
True	0.44 0.85 0.58 2987
accuracy	0.99
macro avg	0.72 0.92 0.79 364539
weighted avg	0.99 0.99 0.99 364539